FAQS

#	Question	Answer
1	What is the hackathon website?	https://lablab.ai/ai-hackathons/amd-developer-hackathon-act-ii
2	I did not get AMD GPU Cloud credits.	You do not need them for the hackathon, you will get a Jupyter Notebook cloud instance made available to you after hackathon start with periods of cloud access.
3	How do I access hackathon compute?	"https://notebooks.amd.com/hackathon" (not "https://notebooks.amd.com") You will be directed to Sign into AMD Developer Program and provided access to your team's compute resource. These are available on hackathon start.
4	Where can I save things for persistent storage for the duration of the hackathon	/workspace
4	How much persistent storage does a team get?	25 GB of persistent store. Please note that it takes a few minutes to flush your local store to persistent store when you request pod deletion.
5	How much compute time does my team get	Each time will get a fixed number of hours over a 24-hour period. The time duration and reset interval (24 hours) will initially be longer but closer to hackathon termination, it will be shorter, to ensure everyone has access to the machines without a long wait.
6	How can one quota consumption when not actually working?	Turn off your session, Pod. Everything in your persistent storage will remain intact. Once a teammate restarts your session, your compute time will start getting consumed. When multiple teammates use a session concurrently, there is not extra compute change. It is no different from one or more persons using a single server machine.
7	I have not got my Fireworks Credits	If you signed up after July 2, you will be allocated credits on July 8.

8.	When is the event submission deadline?	CET 6 pm July 11, 2026
9	When will be results become available?	Approximately a month for the Unicorn track. The other tracks have a leaderboard. All results will be announced together.
10	How do I join a team?	On the LabLab.ai Discord under AMD Hackathon Act-2 there is a sub-channel to request teammates. https://discordapp.com/channels/877056448956346408/1488590902661222600
11	Which models may one use Fireworks AI?	MiniMax and Kimi K **MiniMax** is a series of advanced AI models designed for multimodal understanding, coding, and agentic reasoning, with capabilities spanning text, audio, image, video, and music. MiniMax M3 - Coding & Agentic Frontier, 1M Context, Multimodal | MiniMax **Kimi K** series, from **Moonshot AI**, is a line of open-source large language models designed for **agentic intelligence, coding, and multimodal tasks**. The models use a **Mixture-of-Experts (MoE) architecture**, allowing efficient activation of a subset of experts per token, which enables large-scale reasoning without excessive computational cost
12	How much memory does the GPU on the hackathon compute instance have?	**About 48 GB.** Remember that things like KV-Cache need space too and other infrastructure, so your model needs to be smaller. Please experiment launching with any model you are interested in.
13	How do I launch a model serving instance with vllm?	In an terminal window in your Jupyter pod launch: > **vllm serve Qwen/Qwen2-7B-Instruct --port 8000 --gpu-memory-utilization 0.3**

14	How do I fine-tune a model that needs more training time than my compute window.	Please use checkpointing to save your training state in persistent store. One or two checkpoints should suffice to allow you to pick up where you left off.
		
		Saving
		checkpoint = {
		"epoch": epoch,
		"global_step": step,
		"model_state_dict": model.state_dict(),
		"optimizer_state_dict": optimizer.state_dict(),
		"loss": loss,
		}
		
		torch.save(checkpoint, "checkpoint.pt")
		
		Resuming:
		checkpoint = torch.load("checkpoint.pt")
		
		model.load_state_dict(checkpoint["model_state_dict"])
		optimizer.load_state_dict(checkpoint["optimizer_state_dict"])
		
		start_epoch = checkpoint["epoch"] + 1
		step = checkpoint["global_step"]

15	How do I access Track1 Leaderboard?	https://lablab.ai/ai-hackathons/amd-developer-hackathon-act-ii/live?track=1#amd-leaderboard
16	How do I access Track 2 leaderboard?	https://lablab.ai/ai-hackathons/amd-developer-hackathon-act-ii/live?track=2#amd-leaderboard
17.	Why did I get locked out of my AMD Developer Account while using the AMD AI Academy Jupyter Notebook? How do I get unlocked?	The ai-academy environment is to run notebooks cell-by-cell. If users experiment with new commands, the detection routines block the user. For experiments, they can explore AMD Dev Cloud or our hackathon platform for lablab hackathon. Please send us a support request. Unlocking could take anywhere between 1-3 days, 3 if spanning a weekend.

18	What does N hours in 24 hours mean with respect to compute time?	It means your compute instance will be given resources on our GPU cluster for a maximum of N hours in a 24 hour window. You can use those N hours in one stretch or in multiple steps. To use in steps you need to “stop/shutdown” your instance and “restart” it later. Anyone on your team can shutdown or restart an instance, it is a shared instance. When you restart, if you are using a virtual environment, you will need to use pip install to install any libraries you want in your kernel. Occasionally, under low load we may reset the window to something shorter so you can hop on sooner or even give more hours. Typically close to the end of the hackathon when lots of people are seeking resources this does not happen!
19	Do I need approval for my idea for the hackathon?	No! Welcome! Mentors are a resource, not a gate.
20	How will I get my $50 Fireworks AI credits?	You will receive an email with a code, one per team.
21	Do I need to maintain a live API endpoint for the judges for evaluation?	Absolutely not, it will just be too expensive. If you want you could set up a HuggingFace deployment of your application, which uses a serverless setup to avoid usage credits when not actually in use. For the purposes of hackathon submission you only need to create a demo video of your application, a presentation deck (about 5 slides) and provide your GitHub url.
22	What is the motivation behind Track-1?	Think of enterprises wanting to control their AI spend. Not all their tasks need top-of-the-line proprietary models at premium cost. But they also want a good user experience, not having each user determine which model to use and where. hey may host a bunch of models, typically open source, RAG, and may even be fine-tuned models for their in-house use and delegate to a premium model if necessary. That is where a smart router fits in. When we say tokens expended on local models count as ZERO, we basically are encouraging you to run as many

		local models as you need, including the smart router, so you need to make as few external API calls to the Fireworks AI endpoint.